

Benchmarking deep learning against matched filtering for single-event graviton-dispersion tests on LIGO data

Ram Chand*

Department of Natural Sciences,
The Begum Nusrat Bhutto Women University,
Sukkur, Sindh 65170, Pakistan

September 19, 2026

Abstract

General relativity predicts a massless graviton; a broad class of theories beyond it gives the graviton a small rest mass m_g , which would make gravitational waves *dispersive* and imprint a calculable, frequency-dependent phase on a binary-black-hole inspiral [1, 2]. Measuring that phase tests whether the graviton has mass. We benchmark a deep-learning discriminator against the optimal matched filter on a *single* gravitational-wave event—a demanding regime, since the strongest current bound, $\lambda_g = h/(m_g c) > 9.8 \times 10^{13}$ km, is set only by combining the event catalogue [3], and whether a learned model can beat matched filtering on one event has not been tested on observational strain. We inject general-relativistic and massive-graviton inspirals—source parameters from the GWTC-4.0 catalogue, dispersion phase to leading order in $1/\lambda_g^2$ —into LIGO H1 strain near GW150914 at controlled signal-to-noise ratios $\rho_{\text{opt}} \in \{20, 50\}$, scan the graviton Compton wavelength λ_g , and on identical data compare two discriminators: the *oracle* matched filter, given the true source parameters, and a compact (1.7×10^5 -parameter) convolutional-transformer classifier, given only whitened strain. On observed strain both reach only $\text{AUC} = 0.45\text{--}0.50$ —indistinguishable from each other and from chance at every grid cell—whereas on noise-free signals both separate the classes perfectly. A noise ladder localises the cause: in stationary Gaussian noise of the measured power spectrum the classes stay separable at strong dispersion yet collapse to chance on real strain—so non-Gaussian “glitches” set the single-event limit—while near the observational bound the inter-class difference is itself below the single-event noise. For the model, training and noise studied, deep learning gives no per-event advantage over the Gaussian-optimal matched filter; since that filter is optimal only under the Gaussian likelihood, a method that models the non-Gaussian noise remains a route to improvement we do not exclude. We provide this as a validated, openly reproducible benchmark against which any future claim of a learned single-event advantage can be measured.

Keywords: gravitational waves, massive graviton, tests of general relativity, deep learning, matched filtering, LIGO

1 Introduction

Gravitational waves—propagating ripples in spacetime—are produced when two black holes orbit one another, lose orbital energy to that radiation, and spiral inward until they merge. The signal is a brief “chirp” that sweeps upward in frequency over the final fraction of a second and is recorded on Earth by the kilometre-scale laser interferometers of LIGO and its partner observatories Virgo and KAGRA [4, 5], which sense the minute stretching and squeezing of

*Electronic address: ram.chand2k11@yahoo.com

space the passing wave produces—the *strain*. Since the first such detection, the binary-black-hole merger GW150914, dozens of black-hole mergers have been catalogued, and because each chirp spans a broad band of frequencies it is a sensitive probe of how gravity itself propagates—in particular, of the dispersion relation obeyed by the gravitational field.

In general relativity the graviton—the spin-2 quantum of the gravitational field—is exactly *massless*, so a gravitational wave of every frequency travels at the speed of light and the chirp reaches the detector undistorted. The graviton mass m_g is therefore a sharp parameter with which to test gravity: any nonzero rest mass modifies the gravitational-wave (GW) dispersion relation, so that different frequencies propagate at slightly different speeds and a calculable, frequency-dependent phase shift accumulates on the observed binary-black-hole inspiral waveform [1, 2]. Limits on m_g have historically been set from solar-system dynamics and a range of astrophysical and laboratory probes [6, 7, 8], but the tightest current values come from gravitational-wave inspirals, in which the dispersion accumulates over cosmological propagation distances [9]. From this growing catalogue of black-hole mergers the LIGO–Virgo–KAGRA collaboration has constrained m_g through hierarchical analysis of all events jointly [3]; the third-observing-run combination yields a 90% credible upper limit $m_g c^2 < 1.27 \times 10^{-23}$ eV, or equivalently $\lambda_g = h/(m_g c) > 9.8 \times 10^{13}$ km [3]. This bound is the catalogue-level result; the corresponding constraint from any single event is roughly an order of magnitude weaker, because single-event matched-filter likelihoods are limited by the variance contributed by non-Gaussian noise transients in the LIGO strain [10, 11].

Deep learning has been widely promoted for gravitational-wave analysis, which raises a sharp methodological question: can a learned model extract more information from a *single* event than the optimal matched filter does? In the GW-signal-versus-noise problem, deep classifiers have been shown to recover the sensitivity of matched filtering at substantially lower inference cost, both on synthetic data [12, 13] and on archival LIGO observations [14]; learned models now also perform full Bayesian parameter estimation in near real time [15]. More recent work has stress-tested these methods on realistic detector noise: community mock-data challenges benchmark learned searches against matched filtering on real O3 strain [16], deep residual architectures are trained directly on non-stationary noise [17], and normalising-flow inference has been made statistically exact through importance reweighting [18], together showing that modern architectures can approach matched-filter sensitivity for detection and parameter estimation. The corresponding question for the more specific task of distinguishing two waveform *families* (here, general relativistic versus massive-graviton templates with otherwise identical source parameters) has, to our knowledge, not been examined on observational strain data, and it is this benchmarking question—rather than a new bound on m_g —that we set out to answer. Whether single-event discrimination is limited by the discriminator or by the detector noise is itself an empirical question on real strain; answering it, and identifying *which* feature of the noise is responsible, is part of the benchmark’s value.

We address that question directly. We construct balanced injection sets in which each example is a 4-second segment of H1 strain near the GW150914 epoch, containing a binary black hole inspiral drawn from GWTC-4.0 source parameters and modified by either a Mirshekari–Yunes dispersion phase [2] or by nothing. Each injection is rescaled to a controlled optimal matched-filter signal-to-noise ratio. We then compare two single-event classifiers on the same examples: (i) the oracle matched filter, which has access to the true source parameters (m_1, m_2, D_L) and constructs the GR and massive-graviton templates at those parameters; and (ii) a 1D convolutional plus transformer classifier with no explicit source-parameter input. We find that the two discriminators are statistically indistinguishable and that both perform no better than chance on a single event—a null result we trace to the non-Gaussian detector noise rather than to the choice of method, and which we confirm with two independent controls. Our central contribution is therefore not a new bound on m_g but a validated, openly reproducible benchmark: a controlled injection set, an oracle matched-filter baseline, and a like-for-like learned comparison

on identical data, which together quantify the single-event ceiling and provide a baseline against which future learned methods can be measured.

The paper is organised as follows. Section 2 states the waveform model, the injection procedure, the matched-filter baseline, and the classifier architecture. Section 3 checks the dispersion implementation and shows that the single-event detectability scale matches the catalogue bound to within an order of magnitude. Section 4 reports the per-event sensitivity comparison. Section 5 interprets the comparison in the context of the LIGO–Virgo–KAGRA glitch literature and lists explicit limitations of the present study; Sec. 6 states the conclusions.

2 Methods

2.1 Waveform model

We use the `IMRPhenomD` frequency-domain inspiral–merger–ringdown model of Khan et al. [19, 20], as implemented in `PyCBC` [21], as the general-relativistic baseline waveform $\tilde{h}_{\text{GR}}(f; \theta)$. For the massive-graviton waveform $\tilde{h}_{\text{MG}}(f)$ we add to the GR phase the Mirshekari–Yunes modification [2],

$$\delta\Psi_{\text{MG}}(f) = -\frac{\pi c D_0}{\lambda_g^2 (1+z) f}, \quad (1)$$

where $\lambda_g = h/(m_g c)$ is the graviton Compton wavelength and D_0 is the graviton-distance integral of Mirshekari–Yunes–Will. Equation (1) describes a *propagation* effect: it depends only on distance, λ_g , redshift and frequency, and is independent of the binary’s chirp mass—the mass dependence cancels at the relative 1PN (ppE $b = -1$) order at which the term enters (Appendix C). The expression is dimensionless, $[c D_0/(\lambda_g^2 f)] = 1$. For the events considered ($z \lesssim 0.2$) the cosmological distance D_0 reduces to D_L at the few-percent level, and we take $D_0 = D_L$ in Eq. (1). The massive-graviton waveform is then

$$\tilde{h}_{\text{MG}}(f) = \tilde{h}_{\text{GR}}(f) \exp[i \delta\Psi_{\text{MG}}(f)]. \quad (2)$$

We emphasise that $\delta\Psi_{\text{MG}} \propto D_L$: when we rescale a template’s distance to a different SNR, we must regenerate the dispersion phase at the new distance, not at a fixed reference.

2.2 Injection construction

We draw component masses (m_1, m_2) with replacement from the 56 GWTC-4.0 [22] BBH events with both source-frame masses above $3 M_\odot$ and network matched-filter SNR above 10 (*cf.* Table 1); the catalogue parameters are fetched directly from the `gwosc` JSON endpoint [23, 24]. For each draw and each target optimal SNR ρ_{opt} , we compute a reference GR waveform at $D_L = 100$ Mpc, evaluate its optimal SNR on the strain PSD, and rescale the luminosity distance so that the injected optimal SNR equals ρ_{opt} exactly.

The noise component is taken from a 1024-second segment of Advanced LIGO [4] H1 strain ending 16 s before the GW150914 [5] epoch, downloaded via `gwpv.timeseries.TimeSeries.fetch_open_data` [24]. The one-sided power spectral density $S_n(f)$ is estimated by a median-Welch periodogram with 4-second Hann windows and 50% overlap, following [25]. Each 4-second noise segment used in an injection is drawn from a uniform random offset in the reservoir.

For the classifier input we whiten each 4-second segment by $1/\sqrt{S_n(f)}$ in frequency domain, then standardise each segment to zero mean and unit variance over the interior 75% (to remove whitening edge artefacts). The matched-filter analysis operates directly on the un-whitened strain through the PSD-weighted inner product.

2.3 Matched-filter baseline

We adopt the optimal-filter conventions of [25, 26]: for two frequency-domain signals $\tilde{a}, \tilde{b}$ and a one-sided PSD S_n ,

$$\langle a|b \rangle = 4 \operatorname{Re} \int_{f_{\text{low}}}^{f_{\text{high}}} \frac{\tilde{a}^*(f) \tilde{b}(f)}{S_n(f)} df, \quad (3)$$

with $f_{\text{low}} = 20 \text{ Hz}$ and $f_{\text{high}} = 1024 \text{ Hz}$ throughout. The optimal SNR of a template h in data d is $\rho(h) = \langle d|h \rangle / \sqrt{\langle h|h \rangle}$. For each injection we know the true source parameters (m_1, m_2, D_L) , and the oracle matched filter constructs both $\tilde{h}_{\text{GR}}$ and $\tilde{h}_{\text{MG}}(\lambda_g)$ at these parameters. The discriminator is

$$s_{\text{MF}} \equiv \rho_{\text{MG}}^2 - \rho_{\text{GR}}^2, \quad (4)$$

which in Gaussian noise is twice the log-likelihood ratio between the massive-graviton and general-relativistic hypotheses [25]. *Under the assumed stationary-Gaussian-noise likelihood* this is the strongest possible matched-filter test of the dispersion hypothesis on a single event: any non-oracle template bank produces a discriminator no more powerful than s_{MF} . In genuinely non-Gaussian data the statistic is no longer optimal—a distinction that turns out to be central to interpreting the results (Sec. 5).

2.4 Classifier

We use a one-dimensional convolutional encoder followed by a single transformer encoder layer (Table 2). The strain enters as a length-16384 vector after whitening; three multi-scale convolutional stages (kernel sizes 7, 31, 63, fused 1×1) reduce the sequence length by a factor of 64, after which a single transformer encoder layer with $d_{\text{model}} = 32$, two heads, and feedforward dimension 64 models temporal correlations. The output is global-mean-pooled and projected through a two-layer MLP head to two class logits. Total trainable parameters: 1.73×10^5 .

The classifier is trained with cross-entropy loss, AdamW ($\eta = 3 \times 10^{-4}$, weight decay 10^{-4}), batch size 32, and a cosine-annealed learning-rate schedule. For each $(\lambda_g, \rho_{\text{opt}})$ cell we generate a balanced set of 800 examples (400 per class), split 85%/15% into training and validation partitions (680 training, 120 validation), and train a separate classifier from scratch for 12 epochs. Reported metrics are always evaluated on the held-out validation partition of the same cell.

2.5 Evaluation

For both classifiers we report the area under the receiver-operator characteristic curve (ROC AUC) on the held-out validation partition. Bootstrap 90% confidence intervals on AUC are obtained from 1000 resamples with replacement of the validation set.

3 Implementation validation

We verify the dispersion implementation by computing the normalised match $\mathcal{M}(\tilde{h}_{\text{GR}}, \tilde{h}_{\text{MG}})$, maximised over time and phase, on H1 strain near GW150914 for a GW150914-like template ($m_1 = 36 M_{\odot}, m_2 = 29 M_{\odot}, D_L = 410 \text{ Mpc}$) and a logarithmic scan in λ_g from 10^{13} km to 10^{18} km . The mismatch $1 - \mathcal{M}$ crosses the per-event detectability threshold $1/(2\rho_{\text{opt}}^2)$ at $\lambda_g \approx 5.9 \times 10^{13} \text{ km}$ for this template ($\rho_{\text{opt}} = 35.1$; Fig. 1). This single-event detectability scale is the same order of magnitude as the catalogue-combined LIGO–Virgo–KAGRA bound $\lambda_g > 9.8 \times 10^{13} \text{ km}$ [3], confirming that the phase is implemented with the correct sign, the correct $\propto \lambda_g^{-2}$ scaling, and the correct overall magnitude. The curve is used only as this internal consistency check: a rigorous single-event Bayesian bound is weaker than the catalogue result, which is why a competitive constraint requires combining many events.




Figure 1: Mismatch $1 - \mathcal{M}(\tilde{h}_{\text{GR}}, \tilde{h}_{\text{MG}})$ as a function of λ_g for a GW150914-like template at $D_L = 410$ Mpc, evaluated on the same 1024-second off-source H1 PSD used for the injection study (Sec. 2.2). Dashed line: per-event detectability threshold $1/(2\rho_{\text{opt}}^2)$ for the optimal SNR $\rho_{\text{opt}} = 35.1$ of this template. Dotted line: the LIGO–Virgo–KAGRA catalogue 90% upper limit $\lambda_g > 9.8 \times 10^{13}$ km [3].

4 Results

4.1 Sensitivity on observed strain

Figure 2 shows the validation AUC of both discriminators as a function of λ_g at the two target SNRs. The open squares are the signal-only control: with no strain noise added, the matched-filter statistic s_{MF} separates the GR and massive-graviton classes near-perfectly ($\text{AUC} \gtrsim 0.99$) at every scanned λ_g , because the sign of the log-likelihood ratio is then noise-free. This confirms that the discriminator and the injection pipeline are free of implementation errors.

Once the same injections are embedded in observed H1 strain, the picture changes completely. The oracle matched filter gives $\text{AUC} = 0.45\text{--}0.50$ and the convolutional–transformer classifier $\text{AUC} = 0.46\text{--}0.47$, the two differing by at most ≈ 0.04 at any cell. Every bootstrap 90% confidence interval includes the chance value 0.5, so neither method achieves a statistically significant single-event detection at any $(\lambda_g, \rho_{\text{opt}})$ on the grid, and the two methods are indistinguishable from each other. Increasing the injected SNR from 20 to 50, or decreasing λ_g (which deepens the dispersion signature), does not lift either method off the chance line. The null persists even at $\lambda_g = 10^{12}$ km, where the signal-only control shows the two template families to be all but orthogonal ($\text{AUC} = 1$): a dispersion signature that is trivially separable in the absence of noise is nonetheless undetectable on a single event once observed strain is present.




Figure 2: Single-event AUC for the oracle matched filter (squares) and the convolutional–transformer classifier (circles) as a function of λ_g at two target optimal SNRs. Open squares: matched-filter signal-only control (no strain noise added), near unity ($\text{AUC} \gtrsim 0.99$) because the sign of the log-likelihood ratio is noise-free. Filled symbols: H1 strain near GW150914. Error bars: bootstrap 90% confidence intervals on AUC (1000 resamples). Vertical band: LIGO–Virgo–KAGRA catalogue 90% upper limit on λ_g [3].

4.2 ROC at the most discriminative cell

Figure 3 compares the full ROC curves at the most favourable cell on the grid—the smallest $\lambda_g = 10^{12}$ km at the higher SNR $\rho_{\text{opt}} = 50$, where the signal-only template mismatch is largest. Even here both curves track the chance diagonal, and their bootstrap 90% confidence bands overlap each other and enclose the diagonal along its entire length. The matched filter and the learned classifier therefore deliver the same (null) single-event performance at the point where any dispersion signature should be easiest to detect.

4.3 Positive control: the classifier learns when the noise is removed

The null result of Sec. 4.1 could in principle reflect a network too small, or a training set too sparse, to learn the dispersion signature at all. Figure 4 rules this out. When the identical injections are presented without strain noise (*signal only*), the same compact architecture, trained with the same optimiser and schedule, separates the GR and massive-graviton classes with increasing fidelity as data are added: at $N = 400$ examples per class the noise-free validation AUC rises from chance to ≈ 1.0 within a few epochs [Fig. 4(a)], and the best noise-free AUC grows steadily with training-set size—0.73, 0.94, 1.00 at $N = 100, 200, 400$ [Fig. 4(b)]. On H1 strain, by contrast, the AUC of the same network never departs from chance, at any epoch or any N .

The network therefore has ample capacity to detect the inter-class difference; what defeats it on strain is the detector noise, not a lack of model capacity or training data. Increasing N from 100 to 400 per class does not lift the strain AUC, showing that single-event sensitivity has already saturated with respect to training-set size in this regime—the classifier analogue of the noise-free matched-filter control of Fig. 2.

4.4 Robustness to architecture, seed, and training-set size

Two further controls test whether the null is specific to the compact model or to the modest training sets used above. At the resolvable cell ($\lambda_g = 10^{13}$ km, $\rho_{\text{opt}} = 20$) we repeat the strain benchmark with (i) a second, established architecture—a 1D residual CNN of the type standard in gravitational-wave detection [12, 17], of comparable size ($\sim 5 \times 10^6$ parameters)—alongside the CNN+Transformer; (ii) three independent random seeds per configuration; and (iii) training sets enlarged to $N = 1000$ and $N = 2000$ per class, five times the size used in the main sweep (Fig. 5). Both architectures reach $\text{AUC} \approx 1$ on noise-free signals, confirming that each can learn the dispersion signature. On H1 strain both remain at chance at every training size and seed: the mean validation AUC is 0.51 ± 0.03 (CNN+Transformer) and 0.52 ± 0.02 (ResNet) at $N = 1000$, and 0.50 ± 0.02 and 0.53 ± 0.01 at $N = 2000$. Neither a second, established architecture nor a five-fold increase in training data shifts the single-event strain AUC away from chance—confirming that the null is a property of the data rather than of model capacity, initialisation, or training-set size.

4.5 A noise ladder isolates the cause of the null

The signal-only control shows that the classes are separable in the absence of noise, but not *which* property of the observational data destroys that separability. We localise it by interposing a third noise condition between the two extremes: stationary, colored Gaussian noise whose one-sided power spectrum equals the measured H1 PSD. On identical injections we evaluate the oracle matched filter under three conditions—signal only, Gaussian colored noise, and observational H1 strain—at $\rho_{\text{opt}} = 50$ across the λ_g grid (Fig. 6). The colored-Gaussian generator is normalised so that a unit-norm template projected onto pure noise has unit variance in the inner product of Eq. (3) (measured $\text{Var} = 0.99$); its power therefore matches the PSD the matched filter assumes.

The ladder resolves the null into two regimes. At strong dispersion ($\lambda_g \leq 10^{13}$ km) the two waveform families are nearly orthogonal and their difference carries a large signal-to-noise ratio [$\rho_{\text{diff}} = \rho_{\text{opt}} \sqrt{2(1 - \mathcal{M})} \approx 66$ at $\lambda_g = 10^{12}$ km]: in Gaussian noise the oracle statistic separates them perfectly ($\text{AUC} = 1.00$), yet on real strain it falls to chance ($\text{AUC} = 0.50$). The loss of sensitivity there is caused specifically by the non-Gaussian, non-stationary content of the strain—the glitches of Fig. 8—and not by the signal, the SNR, or a Gaussian noise floor. At weak dispersion ($\lambda_g \geq 10^{14}$ km), near and above the observational bound, the inter-class difference is intrinsically small and the Gaussian-noise AUC itself falls—to 0.56 at $\lambda_g = 10^{14}$ km and to chance (≈ 0.49) by $\lambda_g = 10^{15}$ km—so there the single-event difference-SNR, not the glitch content, is the binding limit. The physically-relevant regime is thus doubly constrained, and the counter-intuitive observation that even the oracle statistic is at chance where the templates are nearly orthogonal is explained: the orthogonality is realised only in the absence of glitches.

To check that the null is not an artefact of one atypical, glitch-rich interval, we repeat the real-strain benchmark on four independent, non-overlapping ~ 4 -minute windows of the O1 H1 reservoir, each with its own measured PSD (Fig. 7). The oracle matched-filter AUC is 0.51 ± 0.02 at $\lambda_g = 10^{12}$ km and 0.51 ± 0.02 at $\lambda_g = 10^{14}$ km across the four windows—chance in every case—so the single-event null is stable against the choice of noise stretch within the reservoir. A wider survey across observing runs and detectors is a natural extension that the released analysis code supports directly.

5 Discussion

The per-event AUC of the oracle matched filter on H1 strain hovers near 0.45–0.50 across the entire scanned $(\lambda_g, \rho_{\text{opt}})$ grid. We have ruled out the possibility that this reflects an implementation defect: on signal-only data the oracle matched filter separates the two waveform families near-perfectly ($\text{AUC} \gtrsim 0.99$) at every scanned λ_g (open symbols, Fig. 2), and the per-event detectability threshold computed from Eq. (3) places the single-event λ_g detectability scale at the same order of magnitude as the catalogue bound (Sec. 3). What dominates the per-event discriminator on strain is the variance contributed by non-Gaussian features in the data — blip glitches [10, 27], harmonic lines, and short-duration transients — whose energy partially projects onto both the GR and the dispersed templates, producing correlated fluctuations in ρ_{GR}^2 and ρ_{MG}^2 that do not cancel in s_{MF} because the two templates have slightly different time–frequency support.

Figure 8 makes this scale separation explicit. Panel (a) shows the noise-free whitened GR and massive-graviton templates for a GW150914-like source at $\rho_{\text{opt}} = 20$, together with their difference $h_{\text{MG}} - h_{\text{GR}}$ on the same vertical scale. Even at $\lambda_g = 10^{14} \text{ km}$ —comparable to the LVK bound—the entire inter-class signal that any single-event discriminator must detect peaks at only ≈ 0.02 noise standard deviations. Panel (b) shows a 4-second segment of off-source H1 strain with the identical injection added, whitened on the same scale: the data are pervasively non-Gaussian, the loudest localised transient reaching $\approx 8\sigma$ —orders of magnitude above the inter-class difference. It is this excess, and not any deficiency of the discriminator, that drives the single-event AUC to chance.

The convolutional–transformer classifier learns no representation that appreciably reduces this glitch-induced variance on the per-event task, at the training-set scales and SNRs studied here. It is important to state precisely what this does and does not establish. The oracle matched filter is optimal only under the assumed stationary-Gaussian-noise likelihood, so its failure on real strain reflects the breakdown of that assumption rather than a fundamental single-event information bound. The noise ladder (Fig. 6) makes the point concrete: the same statistic that is defeated by glitches separates the classes perfectly in Gaussian noise of the same power spectrum. A discriminator that explicitly modelled the non-Gaussian, non-stationary noise—rather than assuming the Gaussian likelihood the matched filter is built on—could in principle exceed s_{MF} on real data, and learning such a noise model is one plausible route by which a per-event advantage might yet emerge; our null does not exclude it. What the result does establish is narrower but firm: for the two architectures, training protocol, noise sample and SNRs studied here—and across three seeds and training sets up to $N = 2000$ per class (Fig. 5)—the learned discriminators match but do not beat the Gaussian-optimal matched filter, and neither is useful per event at GWTC-4.0 SNRs. This is consistent with the LIGO–Virgo–KAGRA practice of obtaining graviton-mass constraints by combining many events.

Limitations. (i) The noise reservoir is restricted to a single ~ 17 -minute segment of O1 H1 strain near GW150914. Within that reservoir the null is stable across independent noise windows (Fig. 7), but other observing runs and detectors carry different glitch populations, and the sensitivity numbers are not guaranteed to generalise across runs without the wider survey the released code enables. (ii) The matched filter is taken in oracle form (templates at the true source parameters). A practical bank-template search [21] or stochastic-sampling parameter estimation [28] would exhibit modestly worse sensitivity than the oracle, so the classifier-versus-MF comparison reported here is conservative against the classifier. (iii) We use leading-order Mirshekari–Yunes dispersion only; subleading effects on merger and ringdown morphology [29] are not modelled. (iv) The training sets are small by GW deep-learning standards [13, 12]; the training-set-size study of Fig. 4(b) and the ablation of Fig. 5, which enlarges N to 2000 per class (five times the main sweep), both show the single-event strain AUC already saturated

at chance, because the per-event variance is set by the strain rather than by the training-set size. We cannot formally exclude modest gains from still larger datasets or a wider sampling of the noise, but they appear unlikely in this regime. (v) The learned discriminators are two of many possible choices; we tested an established 1D ResNet alongside the CNN+Transformer and found the same null (Fig. 5), but specialised models trained on glitch-rich data—transient-aware or attention-based architectures, or discriminators that explicitly learn the non-Gaussian noise likelihood—may give different results, and post-hoc attribution methods [30] could localise which time–frequency features a successful discriminator exploits. We do not claim our architectures are optimal.

6 Conclusion

We set out to test, on a *single* gravitational-wave event, the general-relativistic prediction that the graviton is massless, and to ask whether a learned discriminator can detect a massive-graviton dispersion signature more sensitively than the optimal matched filter. Using balanced injections of general-relativistic and Mirshekari–Yunes–dispersed inspirals—source parameters drawn from the GWTC-4.0 catalogue, dispersion phase implemented to leading order in $1/\lambda_g^2$ —into H1 strain near GW150914, we evaluated a compact convolutional–transformer classifier and the oracle matched filter on identical examples. The two methods are statistically indistinguishable at every $(\lambda_g, \rho_{\text{opt}})$ cell scanned, each saturating in the range $\text{AUC} = 0.45\text{--}0.50$ with confidence intervals that include chance. A control confirms this is a property of the data rather than an implementation defect: on noise-free signals the same matched filter separates the two waveform families near-perfectly ($\text{AUC} \gtrsim 0.99$) at every scanned λ_g , and the single-event detectability scale matches the catalogue-combined bound in order of magnitude [3].

In the physically-relevant regime the single-event sensitivity is therefore set by non-Gaussian detector noise—blip glitches and short-duration transients (Figs. 8, 6)—rather than by the choice between the two discriminators studied: for the two architectures, training protocol and the SNRs typical of the current catalogue, the learned representations match but do not beat the Gaussian-optimal matched filter, and neither offers a shortcut around the noise floor that obliges LIGO–Virgo–KAGRA to combine many events for a useful graviton-mass constraint [3]. This is a null methodological result, and it delimits where future gains must originate. A per-event advantage for deep learning, if one exists, would most plausibly emerge from representations trained explicitly to be robust against glitches, from models that learn to combine many events coherently rather than scoring each in isolation, or from the substantially higher single-event SNRs anticipated with detector upgrades and next-generation observatories. The framework introduced here—open injections, an oracle matched-filter baseline, and a like-for-like classifier comparison on identical data—provides a reusable benchmark against which any future claim of a learned single-event advantage can be tested.

Data and code availability

All code, the curated GWTC-4.0 event list, the result data and figures, and the bibliography are available at <https://github.com/ramc77/gw-graviton-ml-benchmark>. The GW strain data are public via the LIGO Open Science Center, <https://gwosc.org>.

Acknowledgments

The author acknowledges the LIGO Open Science Center for the H1 strain data, the developers of `gwpv` and `PyCBC` for the GW analysis tooling used here, and the Begum Nusrat Bhutto Women University, Sukkur, Pakistan, for institutional support.

References

- [1] Clifford M. Will. Bounding the mass of the graviton using gravitational-wave observations of inspiralling compact binaries. *Physical Review D*, 57(4):2061–2068, 1998.
- [2] Saeed Mirshekari, Nicolás Yunes, and Clifford M. Will. Constraining Lorentz-violating, modified dispersion relations with gravitational waves. *Physical Review D*, 85:024041, 2012.
- [3] R. Abbott et al. Tests of general relativity with gwtc-3. *arXiv preprint*, 2021.
- [4] J. Aasi et al. Advanced LIGO. *Classical and Quantum Gravity*, 32:074001, 2015.
- [5] B. P. Abbott et al. Observation of gravitational waves from a binary black hole merger. *Physical Review Letters*, 116:061102, 2016.
- [6] Alfred S. Goldhaber and Michael Martin Nieto. Photon and graviton mass limits. *Reviews of Modern Physics*, 82:939–979, 2010.
- [7] Claudia de Rham, J. Tate Deskins, Andrew J. Tolley, and Shuang-Yong Zhou. Graviton mass bounds. *Reviews of Modern Physics*, 89:025004, 2017.
- [8] Clifford M. Will. The confrontation between general relativity and experiment. *Living Reviews in Relativity*, 17:4, 2014.
- [9] B. P. Abbott et al. Tests of general relativity with GW150914. *Physical Review Letters*, 116:221101, 2016.
- [10] Miriam Cabero et al. Blip glitches in Advanced LIGO data. *Classical and Quantum Gravity*, 36:155010, 2019.
- [11] Alexander H. Nitz, Thomas Dent, Tito Dal Canton, Stephen Fairhurst, and Duncan A. Brown. Detecting binary compact-object mergers with gravitational waves: understanding and improving the sensitivity of the PyCBC search. *The Astrophysical Journal*, 849:118, 2017.
- [12] Hunter Gabbard, Michael Williams, Fergus Hayes, and Chris Messenger. Matching matched filtering with deep networks for gravitational-wave astronomy. *Physical Review Letters*, 120:141103, 2018.
- [13] Daniel George and E. A. Huerta. Deep learning for real-time gravitational wave detection and parameter estimation: Results with advanced LIGO data. *Physics Letters B*, 778:64–70, 2018.
- [14] Elena Cuoco, Jade Powell, Marco Cavaglià, et al. Enhancing gravitational-wave science with machine learning. *Machine Learning: Science and Technology*, 2:011002, 2021.
- [15] Maximilian Dax, Stephen R. Green, Jonathan Gair, et al. Real-time gravitational wave science with neural posterior estimation. *Physical Review Letters*, 127:241103, 2021.
- [16] Marlin B. Schäfer, Ondřej Zelenka, Alexander H. Nitz, et al. MLGWSC-1: The first machine learning gravitational-wave search mock data challenge. *Physical Review D*, 107:023021, 2023.
- [17] Paraskevi Nousi, Alexandra E. Koloniari, Nikolaos Passalis, Panagiotis Iosif, Nikolaos Stergioulas, and Anastasios Tefas. Deep residual networks for gravitational wave detection. *Physical Review D*, 108:024022, 2023.

- [18] Maximilian Dax, Stephen R. Green, Jonathan Gair, et al. Neural importance sampling for rapid and reliable gravitational-wave inference. *Physical Review Letters*, 130:171403, 2023.
- [19] Sebastian Khan, Sascha Husa, Mark Hannam, Frank Ohme, Michael Pürrer, Xisco Jiménez Forteza, and Alejandro Bohé. Frequency-domain gravitational waves from non-precessing black-hole binaries. II. A phenomenological model for the advanced detector era. *Physical Review D*, 93:044007, 2016.
- [20] Sascha Husa, Sebastian Khan, Mark Hannam, et al. Frequency-domain gravitational waves from nonprecessing black-hole binaries. I. new numerical waveforms and anatomy of the signal. *Physical Review D*, 93:044006, 2016.
- [21] Samantha A. Usman et al. The PyCBC search for gravitational waves from compact binary coalescence. *Classical and Quantum Gravity*, 33:215004, 2016.
- [22] LIGO Scientific Collaboration, Virgo and KAGRA Collaborations. Gwtc-4.0: The fourth gravitational-wave transient catalog. *arXiv preprint*, 2023.
- [23] R. Abbott et al. Open data from the first and second observing runs of advanced LIGO and advanced virgo. *SoftwareX*, 13:100658, 2021.
- [24] Duncan M. Macleod et al. GWpy: A python package for gravitational-wave astrophysics. *SoftwareX*, 13:100657, 2021.
- [25] Bruce Allen, Warren G. Anderson, Patrick R. Brady, Duncan A. Brown, and Jolien D. E. Creighton. FINDCHIRP: An algorithm for detection of gravitational waves from inspiralling compact binaries. *Physical Review D*, 85:122006, 2012.
- [26] Curt Cutler and Éanna E. Flanagan. Gravitational waves from merging compact binaries: How accurately can one extract the binary’s parameters from the inspiral waveform? *Physical Review D*, 49:2658–2697, 1994.
- [27] Michael Zevin et al. Gravity Spy: Integrating advanced LIGO detector characterization, machine learning, and citizen science. *Classical and Quantum Gravity*, 34:064003, 2017.
- [28] Gregory Ashton, Moritz Hübner, Paul D. Lasky, et al. BILBY: A user-friendly bayesian inference library for gravitational-wave astronomy. *The Astrophysical Journal Supplement Series*, 241(2):27, 2019.
- [29] Nicolás Yunes and Frans Pretorius. Fundamental theoretical bias in gravitational-wave astrophysics and the parameterised post-Einsteinian framework. *Physical Review D*, 80:122003, 2009.
- [30] Scott M. Lundberg and Su-In Lee. A unified approach to interpreting model predictions. *Advances in Neural Information Processing Systems*, 30, 2017.

A Curated GWTC-4.0 event list

B Classifier architecture

C Massive-graviton dispersion phase

The dispersion relation for a massive graviton of rest mass m_g gives a group velocity $v_g/c = \sqrt{1 - (m_g c^2/E)^2}$ with $E = hf$. Writing $m_g c^2 = hc/\lambda_g$ and expanding to leading order in $(c/\lambda_g f)^2$, a graviton of frequency f accrues, over a propagation distance D_0 , a delay $\Delta t(f) =$

Table 1: Source-frame masses, luminosity distance, and network matched-filter SNR for the GWTC-4.0 BBH events used as the parameter pool for our injection campaign ($m_1, m_2 > 3 M_\odot$; network SNR > 10). Top 12 by SNR shown; full list (56 events) in `gwtc4_bbh_events.csv` in the accompanying code repository. Values from the `gwosc.api` event JSON endpoint.

Event	$m_1 (M_\odot)$	$m_2 (M_\odot)$	D_L (Mpc)	ρ
GW230814.230901	33.7	28.2	290	43.0
GW231226.101520	40.2	35.1	1180	34.7
GW230627.015337	8.4	5.7	305	28.7
GW231028.153006	94.0	59.0	4200	22.4
GW231206.233901	37.6	28.5	1490	21.9
GW231123.135430	137.0	101.0	2200	21.8
GW230927.153832	21.7	16.6	1190	20.3
GW230914.111401	60.0	37.0	2700	17.2
GW230919.215712	27.3	21.4	1330	16.8
GW230628.231200	32.5	27.0	2260	16.4
GW231102.071736	61.0	42.0	3800	15.6
GW240104.164932	42.3	32.2	1910	14.8

Table 2: Layers and parameter counts in the compact convolutional plus transformer classifier. Total trainable parameters: 1.73×10^5 .

Layer	Output / parameters
Input strain	$(B, 1, 16384)$
Multi-scale conv stage 1 (kernels 7, 31, 63)	$(B, 12, 4096)$
Multi-scale conv stage 2	$(B, 24, 1024)$
Multi-scale conv stage 3	$(B, 48, 256)$
1×1 projection to $d_{\text{model}} = 32$	$(B, 32, 256)$
Sinusoidal positional encoding	—
Transformer encoder layer (2 heads, FF = 64)	$(B, 256, 32)$
Global mean pool	$(B, 32)$
MLP head ($32 \rightarrow 64 \rightarrow 2$)	$(B, 2)$

$D_0 c / (2\lambda_g^2 f^2)$ relative to a massless one. The corresponding stationary-phase correction to the frequency-domain waveform is $\delta\Psi_{\text{MG}}(f) = -2\pi f \Delta t(f)$, with the pieces constant and linear in f reabsorbed into the coalescence time and phase. Restoring the cosmological $(1+z)$ factor [1, 2],

$$\delta\Psi_{\text{MG}}(f) = -\frac{\pi c D_0}{\lambda_g^2 (1+z) f}, \quad (5)$$

with D_0 in meters, λ_g in meters, f in hertz, and $c = 2.998 \times 10^8$ m/s. The expression is dimensionless and contains no chirp mass: graviton dispersion is a propagation effect, and the binary-mass dependence cancels between the ppE amplitude and the $(\pi\mathcal{M}_c f)^{-1}$ factor at the relative 1PN ($b = -1$) order at which the term enters.




Figure 3: ROC at the cell with the largest signal-only mismatch in the scanned grid. Shaded bands: bootstrap 90% confidence regions.




Figure 4: Classifier positive control at a resolvable-but-buried cell ($\lambda_g = 10^{13}$ km, $\rho_{\text{opt}} = 20$). (a) Validation AUC versus training epoch at $N = 400$ examples per class: noise-free (signal-only) injections (circles) versus H1 strain near GW150914 (squares). (b) Best validation AUC versus training-set size N per class. Noise-free, the network’s AUC climbs toward unity as N grows; on strain it remains at chance (dotted line). The contrast establishes that the null result is set by the detector noise, not by network capacity or training-set size.




Figure 5: Architecture, seed, and training-size ablation at $\lambda_g = 10^{13}$ km, $\rho_{\text{opt}} = 20$. Best validation AUC versus training-set size for the compact CNN+Transformer (blue) and an established 1D ResNet baseline (orange) on H1 strain (circles, mean over three seeds at $N = 1000$ and 2000 and two seeds at $N = 400$; error bars $\pm 1\sigma$), with the noise-free capability control for each architecture (stars, AUC ≈ 1). Both architectures remain at the chance line (dotted) on strain for all training sizes up to $N = 2000$ per class, while both separate the classes perfectly in the absence of noise.




Figure 6: Noise ladder: oracle matched-filter AUC versus λ_g at $\rho_{\text{opt}} = 50$ under three noise conditions on identical injections—signal only (open circles), stationary colored Gaussian noise with the measured H1 PSD (squares), and observational H1 strain (filled circles); 90% bootstrap intervals. At $\lambda_g \leq 10^{13}$ km the classes are trivially separable in Gaussian noise yet at chance on real strain, isolating non-Gaussian transients as the cause; at $\lambda_g \geq 10^{14}$ km the Gaussian-noise AUC itself falls toward chance as the inter-class difference becomes intrinsically too small.




Figure 7: Robustness of the null to the noise reservoir. Real-strain oracle matched-filter AUC across four independent time windows of the O1 H1 reservoir, at $\lambda_g = 10^{12}$ km (strong dispersion) and $\lambda_g = 10^{14}$ km (near the LVK bound), $\rho_{\text{opt}} = 50$, $N = 400$ per class; 90% bootstrap intervals. Every window sits at the chance line.




Figure 8: Why single-event discrimination is noise-limited. (a) Noise-free whitened GR template, massive-graviton template ($\lambda_g = 10^{14}$ km), and their difference (offset below, same vertical scale) for a GW150914-like source at $\rho_{\text{opt}} = 20$; the inter-class difference peaks at $\approx 0.02\sigma$. (b) A 4-s off-source H1 strain segment near GW150914 with the identical GR injection added, whitened on the same noise scale; the loudest localised noise transient (marked) and the pervasively non-Gaussian baseline dwarf the inter-class difference of panel (a). Both panels are in common units of detector-noise standard deviations.